

Dell Crowbar Software Framework User's Guide

Version 1.6

July 24, 2013

Contents

Contents	2
Figures	4
Tables.....	4
Notes, Cautions, and Warnings	5
Trademarks.....	5
Introduction	6
Concepts.....	6
Opscode Chef Server.....	6
New for Version 1.6.....	6
The Crowbar Framework	7
Architecture.....	7
System End State	7
Node Provisioning	7
NTP.....	8
DNS	8
Nagios®	8
Ganglia.....	8
Logging.....	8
Network Setup	9
Default Networks.....	9
<i>IP Addressing.....</i>	<i>10</i>
Crowbar Operation	10
<i>Barclamp Life Cycle.....</i>	<i>10</i>
<i>Crowbar State Machine.....</i>	<i>11</i>
<i>Discovering Node.....</i>	<i>12</i>
<i>Allocating Node.....</i>	<i>12</i>
<i>Hardware Install.....</i>	<i>12</i>
<i>Base OS Install</i>	<i>12</i>
<i>Ready for Role</i>	<i>13</i>
<i>Hardware Update</i>	<i>13</i>
<i>Applying Role.....</i>	<i>13</i>
User Interface/Using Crowbar	13
<i>Nodes</i>	<i>13</i>
<i>Barclamps.....</i>	<i>14</i>
General Layout.....	14
Nodes (System Dashboard)	14
<i>Node Alias and Description.....</i>	<i>14</i>
<i>Node Groups.....</i>	<i>14</i>
<i>Node Details.....</i>	<i>15</i>
<i>Bulk Edit</i>	<i>16</i>
Network.....	16
<i>Switches.....</i>	<i>17</i>
<i>VLANs</i>	<i>17</i>
Barclamps	17

<i>Barclamp List</i>	18
<i>Proposal View/Edit</i>	18
Utilities	19
<i>Exported Items</i>	19
<i>Tempest</i>	20
<i>Swift Dispersion</i>	20
Core Crowbar Barclamps	21
Crowbar Barclamp	21
Deployer Barclamp	22
Provisioner Barclamp	23
Network Barclamp	23
Dell RAID Barclamp	25
<i>Operations</i>	25
<i>Supported RAID Controllers</i>	25
BIOS Barclamp	26
<i>PowerEdge-C Series</i>	26
<i>PowerEdge-R Series</i>	26
<i>Operations</i>	26
<i>Firmware Updates</i>	27
<i>Updating BIOS and BMC Firmware</i>	27
<i>Whitelisting New BIOS Versions</i>	28
Dell Branding Barclamp	28
IPMI Barclamp	29
NTP Barclamp	29
DNS Barclamp	29
Logging Barclamp	30
Nagios® Barclamp	30
Test Barclamp	30
Ganglia Barclamp	30
Git Barclamp	31
Supplemental Material	32
System Verification	32
Managing Barclamps.....	32
<i>Introduction</i>	32
<i>Importing a Barclamp</i>	32
<i>More Information</i>	32
Support	33
Crowbar Support	33
To Learn More	33

Figures

Figure 1: Life Cycle of a Barclamp from Concept to Proposal and Deployment	10
Figure 2: The States a System Progresses through to Fully Provision a System	12

Tables

Table 1: Default Networks	9
Table 2: Network Parameters	9
Table 3: Default Network Addresses	10
Table 4: Changes to proposals based on their current state	11
Table 5: Service URLs	13
Table 6: Deployment States	15
Table 7: Details Panel Buttons	15
Table 8: Buttons on Bottom of Details Panel	16
Table 9: Proposal Status	18
Table 10: Barclamps	21
Table 11: Crowbar Barclamp Parameters	22
Table 12: User Name Key	22
Table 13: Deployer Barclamp Parameters	22
Table 14: BIOS Map Entry Keys	23
Table 15: Provisioner Barclamp Parameters	23
Table 16: Network Configuration Options	24
Table 17: Teaming Sub-Parameters	24
Table 18: Default Networks	24
Table 19: Network Parameters	24
Table 20: Range Map String Key	24
Table 21: Supported RAID Controllers	25
Table 22: RAID Controller Capabilities	26
Table 23: Supported Firmware Updates	27
Table 24: IPMI Barclamp Parameters	29
Table 25: NTP Barclamp Parameters	29
Table 26: DNS Barclamp Parameters	29
Table 27: Logging Barclamp Configuration Parameters	30
Table 28: Nagios [®] Barclamp Parameters	30
Table 29: Ganglia Barclamp Parameters	30
Table 30: Test Barclamp Parameters	30
Table 31: Ganglia Barclamp Parameters	30
Table 32: Git Barclamp Parameters	31

Notes, Cautions, and Warnings



A **NOTE** indicates important information that helps you make better use of your computer



A **CAUTION** indicates potential damage to hardware or loss of data if instructions are not followed.



A **WARNING** indicates a potential for property damage, personal injury, or death.



The **OPSCODE LOGO** indicates additional Opscode Chef Server information.



The **DELL LOGO** indicates additional Dell-specific information.

Trademarks

Information in this document is subject to change without notice.

© 2013 Dell Inc. All rights reserved.

Reproduction of these materials is allowed under the Apache 2 license.

Trademarks used in this text: Dell™, the DELL logo, Nagios®, Ganglia, Opscode Chef™, OpenStack™, Canonical Ubuntu™, VMware™, Dell Precision™, OptiPlex™, Latitude™, PowerEdge™, PowerVault™, PowerConnect™, OpenManage™, EqualLogic™, KACE™, FlexAddress™ and Vostro™ are trademarks of Dell Inc. Intel®, Pentium®, Xeon®, Core™ and Celeron® are registered trademarks of Intel Corporation in the U.S. and other countries. AMD® is a registered trademark and AMD Opteron™, AMD Phenom™, and AMD Sempron™ are trademarks of Advanced Micro Devices, Inc. Microsoft®, Windows®, Windows Server®, MS-DOS® and Windows Vista® are either trademarks or registered trademarks of Microsoft Corporation in the United States and/or other countries. Red Hat Enterprise Linux® and Enterprise Linux® are registered trademarks of Red Hat, Inc. in the United States and/or other countries. Novell® is a registered trademark and SUSE™ is a trademark of Novell Inc. in the United States and other countries. Oracle® is a registered trademark of Oracle Corporation and/or its affiliates. Citrix®, Xen®, XenServer® and XenMotion® are either registered trademarks or trademarks of Citrix Systems, Inc. in the United States and/or other countries. VMware®, Virtual SMP®, vMotion®, vCenter®, and vSphere® are registered trademarks or trademarks of VMware, Inc. in the United States or other countries.

Other trademarks and trade names may be used in this publication to refer to either the entities claiming the marks and names or their products. Dell Inc. disclaims any proprietary interest in trademarks and trade names other than its own.


July 24, 2013

Introduction

This document provides instructions for operating the Dell™ Crowbar software framework. Please refer to additional user's guides for specific products that are deployed by Crowbar, such as OpenStack™ or Apache™ Hadoop™.


Concepts

The purpose of this guide is to explain the user interface of Crowbar. Use the *Crowbar Deployment Guide* for assistance with installing Crowbar and configuring the target systems.

 Concepts beyond the scope of this guide are introduced as needed in notes and references to other documentation

Opscode Chef Server

Crowbar makes extensive use of Opscode Chef Server, <http://opscode.com>. To explain Crowbar actions, you should understand the underlying Chef implementation.

 To use Crowbar, it is not necessary to log into the Chef Server; consequently, use of the Chef UI is not covered in this guide. Supplemental information about Chef is included.

This guide provides this additional Chef information as notes flagged with the Opscode logo.

New for Version 1.6

This section outlines new Crowbar features.

- Dell PowerEdge-C6220 and C8000 support for Compute nodes
- Dell Force10 S4810 support for ToR and aggregation switches
- 10 GbE support on the following platforms:
 - C6220
 - C8000
 - R720
- 3 TB large disk support on the following platforms:
 - C8000
 - R720
- Scalability up to 6 to 120 nodes per cluster

The Crowbar Framework

Architecture

Crowbar provides a modular platform containing the building blocks to provision, monitor, and operate a large-scale cloud deployment. Starting with bare metal installation, Crowbar automates all the required installation and configuration tasks. The core capabilities provided are:

- Hardware configuration — updating and configuring BIOS and BMC boards
- Deployment of base operating system
- Deployment of cloud components
- Providing core network infrastructure services (NTP, DNS, DHCP)
- Monitoring availability and performance of all deployed components

To accomplish this, Crowbar is installed on a dedicated administration node. From this admin node, Crowbar manages the initial discovery and configuration of the other nodes in the system. Each function in the cloud is controlled by a Crowbar component called a barclamp. There are barclamps for Nagios®, Ganglia, NTP, and a variety of other basic infrastructure services. Each barclamp is responsible for all the aspects of the underlying technology required to make it usable. To control the operation of a barclamp, you create a proposal for the barclamp (or may edit one already in place). A proposal comprises several parts:

- Parameters to customize the operation of the barclamp; for example: upstream DNS resolvers.
- List of machines in the deployment that fulfill the different roles in the barclamp.
- Internal system information.

When provisioning a function, you start with a proposal generated by Crowbar. Each core service running on the admin server has a default proposal included as part of the Crowbar installation. You can edit these proposals before installing these services on the admin node.

When a proposal is committed, Crowbar configures the Chef server and other components in the systems (TFTP, DHCP, and so on) to build the setup described in the proposal. Machines in the deployment affected by the proposals have their configuration updated using Chef client-side components. At the end of the process, the function described by the proposal is ready for use.

To allow easy integration into your existing environments, Crowbar allows customization of its barclamps, and additional barclamps can be added. You can disable the default monitoring tools (Nagios® and Ganglia) if you prefer to use your own existing monitoring tools, and internal cloud services can be connected to extant services; for example the cloud's NTP service can be configured to synchronize with existing servers.

Finally, a cloud deployment is dynamic. Machines come and go, break down, or get repurposed. Crowbar's operational model makes sure that machines are hooked into the key infrastructure services. Critical services (for example Nagios® monitoring) are installed automatically on newly provisioned machines by default, and those machines may be easily allocated for use with any additional Crowbar services desired.

System End State

The sections that follow describe the services and capabilities available, assuming the system is installed with defaults. As mentioned above, the Crowbar framework allows for many customizations. This section focuses on the primary use cases for Crowbar, namely integrating all the functions into an existing network environment. Later sections describe more advanced customization options.

Node Provisioning

When a new node is added to the system, the node should be set up to allow PXE booting. Once a machine is powered on, Crowbar uses the PXE boot protocol to manage the provisioning process.

After a system is fully installed by Crowbar, it has the following characteristics:

- BIOS is updated and configured based on the system's usage.

- BMC (baseboard management controller) is configured to allow management and IPMI support.
- Base operating system (OS) is installed.
- Administration access to the OS is configured — IP addressing and SSH keys are installed.
- Nagios® and Ganglia monitoring scripts are installed for the functions deployed on the system.
- Chef-client daemon is configured to maintain the system's state in sync.
- NTP sync client is configured.

Additionally, the system is configured to fulfill the functions that are deployed on it and is added to the appropriate cluster.


Crowbar's network barclamp carries on responsibilities related to L2/L3 management, namely:

- Networking administration.
- Physical NIC configuration — BMC port allocation (teamed or not).
- VLAN configuration on nodes.
- IP address location service, used by the rest of Crowbar. Addresses can be allocated from different pools, meant for different usages; for example: admin network, BMC, storage, and public.

The above functions involve managing information on the admin node and remotely executing operations on the compute nodes as they are provisioned. On each compute node, NICs are defined to match VLANs and appropriate addressing information is configured.

NTP

The admin node runs an NTP server to synchronize time on all the machines in the cluster. Optionally, the NTP server can synchronize with upstream servers, in which case nodes are configured to sync their local time to those servers instead.

 Since all nodes in the cluster rely on the admin node, it is important that the time, date, and time zone are correctly set on this node. Crowbar expects the hardware clock on the admin node to be set to the UTC time zone.

DNS

The admin node runs a DNS server to allow resolution of internal and (optionally) external names. The DNS server can be configured with the following:

- A list of upstream DNS servers to contact.
- A set of static mappings.
- The default domain name suffix.
- Crowbar makes sure that when a new machine is added to the deployment, it has a default entry added to the DNS zone. The default host name is the machine's MAC address prefixed by the letter "d" (for example: d00-a4-b3-c2-d1-e0.yourdomain.com).

Nagios®

Nagios® monitors provisioned services for availability. Each cluster instance is represented as a host group, letting you quickly identify the health of a given instance.

Crowbar installs the Nagios® server on the admin node and configures it to monitor all the nodes in the system. As new nodes are brought online, Crowbar dynamically updates Nagios® to include them.

Ganglia

Ganglia monitors the installed cluster for capacity and performance information, letting you easily gauge the cluster's capacity and check recent activity.

Logging


The admin node serves as a central log repository. Its syslog daemon is configured to accept remote messages and each node is configured to forward all messages there.

Network Setup

Due to the nature of the different software used, the network is set up as flat as possible using a dedicated BMC port and bonded LOMs. Crowbar manages all networks, and comes out of the box preconfigured to allow the initial configuration to come up quickly by predefining the storage, admin, public, and BMC networks.

The Crowbar network configuration can be customized to better map to site-specific networking needs and conventions. These changes include adding additional VLANs, changing VLAN mappings, and teaming NICs.

Please refer to the [Network Barclamp](#) section for specific details.


 Networks for the environment are configured when the Crowbar installation is performed. They cannot be changed without re-installing Crowbar. For specific information on how to configure the network JSON file, please see the Crowbar Deployment Guide.

Default Networks

The default networks are presented in the following table. These defaults can be modified prior to installing Crowbar for your specific environment.

Table 1: Default Networks

Usage	Description	Default reserved VLAN tag	Tagged
Admin/Internal VLAN	Used for administrative functions such as Crowbar node installation, TFTP booting, DHCP assignments, KVM, system logs, backups, and other monitoring. There is only one VLAN set up for this function and it is spanned across the entire network.	100	Not tagged
BMC VLAN	Used for connecting to the BMC of each node.	100	Not tagged
Storage VLAN	Used by the Swift storage system for replication of data between machines, monitoring of data integrity, and other storage-specific functions.	200	802.1q Tagged
Edge/External VLANs	Used for connections to devices external to the Cloud infrastructure; these include externally visible services, such as load balancers and web servers. Use one or many of these networks, dependent on the need to segregate traffic among groups of servers.	300	802.1q Tagged

 The admin and BMC networks are expected to be in the same L2 network.

Each network defined in the system has the following parameters:

Table 2: Network Parameters

Name	Default	Description
vlan	Integer	The VLAN to use on the switch and interfaces for this network
use_vlan	true	A value of true indicates that the VLAN should be applied to the interface. A value of false assumes that the node will receive untagged traffic for this network.
add_bridge	false	Indicates if the network should have a bridge built on top of it. If a bridge is created, it will be named "br0".
subnet	IP Address	The subnet for this network
netmask	Netmask	The netmask for this network


Name	Default	Description
router	IP Address	The default router for this network
broadcast	IP Address	The default broadcast address for this network
ranges	map	This contains a map of strings to start and stop values for the network. This allows for sub-ranges with the network for specific uses, e.g. dhcp, admin, bmc, hosts.

IP Addressing

By default, IP addresses are assigned in the following fashion, using large subnets to support many machines on the production network. The table below shows the networks that are installed when the default network configuration is used. In each network, the first 10 IP addresses are reserved for switches, routers, and firewalls.

Table 3: Default Network Addresses

LAN	VLAN	Network	Subnet	Gateway	802.1q	Bridged
Storage	200	192.168.125.0	255.255.255.0	none	Yes	No
Public	300	192.168.126.0	255.255.255.0	192.168.126.1	Yes	No
Admin	100	192.168.124.0	255.255.255.0	192.168.124.1	No	No
BMC	100	192.168.124.0	255.255.255.0	192.168.124.1	No	No
BMC_Vlan	100	192.168.124.0	255.255.255.0	192.168.124.1	No	No


 Each network's ".1" address is reserved for the network gateway.

Crowbar Operation

Barclamp Life Cycle

Understanding the barclamp life cycle is essential to understanding the Crowbar user interface layout.

Figure 1: Life Cycle of a Barclamp from Concept to Proposal and Deployment




Crowbar progresses all deployments through a fixed lifecycle. It is important to understand this life cycle to use Crowbar.

Figure 1 shows the life cycle of a barclamp within Crowbar. A barclamp defines the capability for a service but cannot be deployed. To deploy a barclamp, you must create a proposal. Once the proposal is created, you must select one or more nodes that will fulfill the roles provided by the barclamp. As discussed in the next sections, you may also edit the proposal's attributes as needed.

Applying the proposal tells Crowbar to deploy the proposal onto the nodes. Nodes return to the Ready state when deployment is completed. At this point, the proposal is considered active.

Once a proposal has become active, you may still edit it and reapply the changes. Crowbar will adjust the deployed proposal. Removing nodes will change the configuration for nodes remaining in the proposal; however, it does *not* perform cleanup actions on the removed node. You may also "deactivate" an applied proposal. This will remove the nodes from the proposal but keeps the configuration data.







If you delete an inactive proposal then your configuration changes will be lost.

 Many barclamps require multiple nodes before they can be deployed. A proposal is not deployed until it has a sufficient number of nodes

At the time a proposal is applied, Crowbar updates the "Run List" of the Crowbar managed node role in Chef.

The following table shows options for changes to proposals based on their current state.

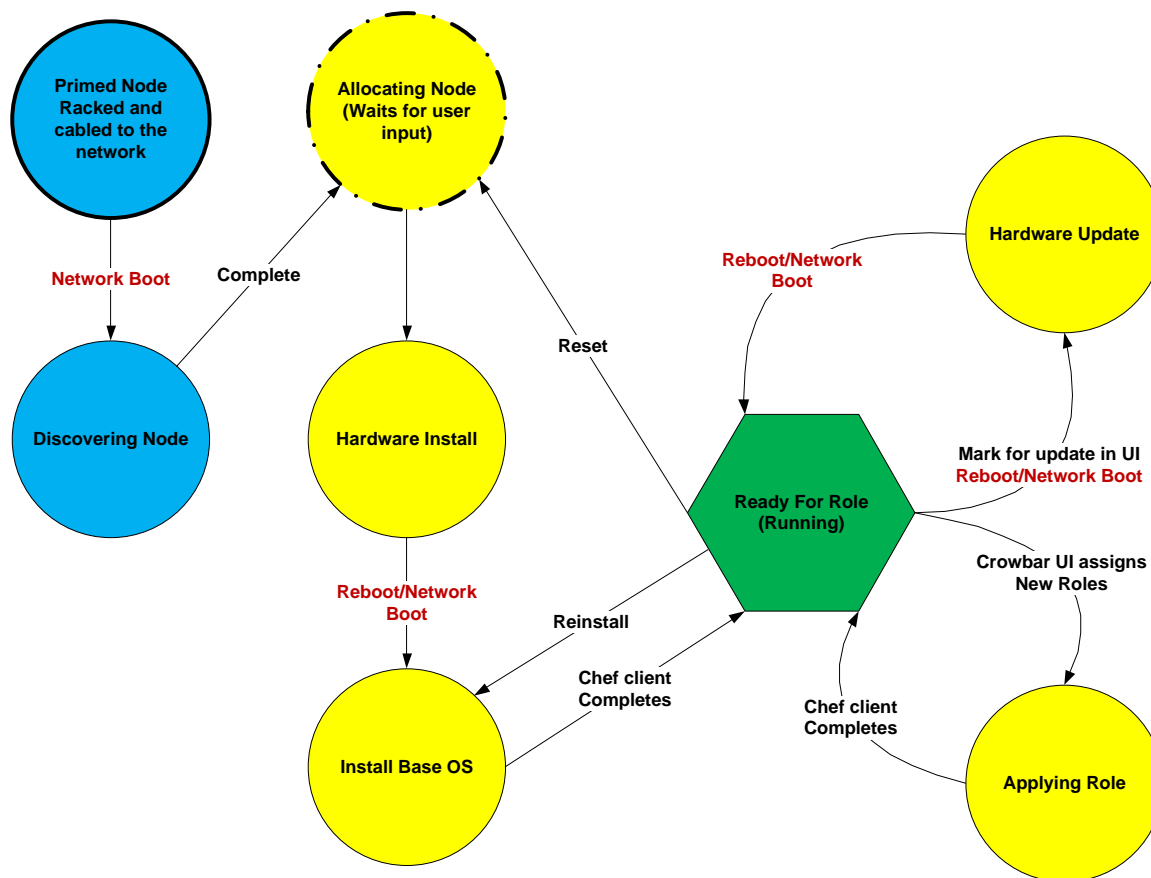
Table 4: Changes to proposals based on their current state

State	Icon	Forwards	Backwards	Comment
No Proposal		Create	None	You must create proposal
User Input		Apply	Delete	
Active		Apply	Deactivate	You may reapply without deactivating a proposal.
Pending	 blinking	None: Proposal Queued	Dequeue	Crowbar is building nodes during this phase.
In Progress	 spinning	None: System Working	None: System Working	
Failed	 blinking	Apply (after fixing issue)	Delete	In some cases, reapplying the proposal is sufficient to fix.

Crowbar State Machine

When Crowbar brings systems up, the systems run through a series of states, each of which performs different actions on the systems. This is controlled by the Provisioner barclamp, and the status of any of the systems is indicated by the icon next to the system on the Dashboard page. You can see a description of the state a system is in by hovering over the status icon.

Figure 2: The States a System Progresses through to Fully Provision a System



Discovering Node

During the Discovering Node state, the node network boots a CentOS LiveCD image to make sure that the node can run some basic components. Once this is complete, the LiveCD image registers with the Admin node to declare its state. The admin node adds the machine to the Chef database, and allocates a BMC and Admin network address for the node. After discovery, no persistent changes have been made to the node or its configuration. The node must be allocated before any permanent changes are made.

Allocating Node

By default the nodes pause after the discovered state. The node awaits allocation either manually through the API (CLI or UI) or by being included in a proposal that is committed to the system. Since the node is not in the ready state, the proposal is queued until the newly allocated node is ready. During the allocation process, Crowbar examines what roles will be applied to the node and uses this information to influence the configuration of the BIOS and RAID systems.

Hardware Install

The LiveCD image reboots forcing another network boot. During the Hardware Install state, the node network boots a CentOS LiveCD image to make sure that the BIOS/BMC/RAID controller are up-to-date. The CentOS LiveCD also updates the BIOS and RAID configurations. Once this is complete, the LiveCD image registers with the Admin node to declare its state. The Admin node updates the Chef database and resets the DHCP entry to installation state.

Base OS Install

The LiveCD image reboots forcing another network boot. The node boots into a network installation of Ubuntu or Red Hat. The installation process reboots the node. Upon reboot and completing the one-time setup script, the system transitions to the Ready for Role state. The first time the Chef

client is run, the node gets a base set of configuration (Nagios® client, Ganglia client, NTP client, DNS resolver). Reinstall will reinstall the base OS on the node, as well as redeploying the roles applied to the node. This effectively uninstalls any custom software that was installed on it and removes all custom configuration that was done to it.

Ready for Role

This is the state that a node spends most of its time in. The node runs the Chef client periodically to ensure that it is up-to-date. It waits for changes to its configuration. If the node reboots, then the node network boots the local boot image again and returns to this state.

Hardware Update

To transition to this state, mark the node as needing hardware updates in the Crowbar UI and then reboot the node. The node network boots and does the same actions as the Hardware Install state, except that it makes no changes to the RAID configuration. Once this is complete, the LiveCD image registers with the Admin node to declare its state. The Admin node updates the Chef database and resets the DHCP entry to installation state. The return state is back to the Ready for Role state. This state allows for the update of the BIOS on a provisioned node without making any other changes to the server.

Applying Role


This is a transient state that represents the running Chef client. It is the time that the node is applying a new configuration before returning to Ready for Role. The transition to this state happens periodically or can be forced by the Admin node.

User Interface/Using Crowbar

Crowbar is delivered as a web application available on the Admin node using HTTP on port 3000. By default, you can access it using <http://192.168.124.10:3000> (see table below). The Crowbar dashboard includes links to the other tools installed, including Ganglia, Nagios®, and Chef.

Table 5: Service URLs

Service	URL	Credentials
SSH	root@192.168.124.10	crowbar
Crowbar UI	http://192.168.124.10:3000/	crowbar / crowbar
Nagios® UI	http://192.168.124.10/nagios3	nagiosadmin / password
Ganglia UI	http://192.168.124.10/ganglia	nagiosadmin / password
Chef UI	http://192.168.124.10:4040/	admin / password

 Crowbar is supported on the following browsers: Firefox 3.6 and later, Internet Explorer 8, Internet Explorer 9, Google Chrome, and Safari 5. HTML5 compatibility and a minimum screen resolution of 1024x768 are recommended.

The IP address 192.168.124.10 is the default address. Replace it with the address assigned to the Admin node.

The Crowbar interface has two primary concepts: nodes and barclamps. All actions are focused on management of these two elements. Before discussing the UI, it's important to understand how they are used by Crowbar.

Nodes

Nodes represent distinct servers in your environment. A server is a single operating system with multiple NICs and HDDs. Each server is identified uniquely by the MAC address of the NIC on the administrative network.



Crowbar nodes map directly to Chef nodes. In fact, all data used in Crowbar is stored entirely in Chef. Chef is the database for Crowbar. Changing a node's data in Chef changes it in Crowbar.

Barclamps

- Barclamps represent modular capabilities that you can install onto none, some, or all of the nodes in your environment.
- Barclamps are activated by generating a proposal for that Barclamp. It is possible to generate multiple proposals for some barclamps, while others may only have a single instance.
- Once a proposal is reviewed, you must activate it before it becomes active in the system.



Barclamps are decomposed in Chef as multiple components: Crowbar data bag entries, cookbooks, recipes, and roles. Our objective is to allow the Chef components used by barclamps to operate in Chef even without Crowbar.

Barclamps have a specific life cycle that is discussed below in more detail as we explore the user interface. Information about using, creating, and extending barclamps is included in the Supplemental Material section.

General Layout

The menu for Crowbar is displayed on the upper right side of the page. Clicking on one of the menu sections causes related content to display on the lower section of the screen.

Alerts or confirmation messages may be displayed between the menu and the page content. Most Crowbar screens automatically update state information so you should not have to refresh the page to get the most current information.



Nodes (System Dashboard)

The Dashboard shows all the nodes in the system and lets you manipulate their power and configuration states.


Node Alias and Description

Users can override the default MAC addressed-based names in the UI from the Node edit page or Bulk Edit pages. When set, the entire UI will display the Node's alias instead of its name. The system name and description are displayed when the mouse hovers over the node alias.

Node Groups

The top of the group box (red in the illustration above) shows the Group Name and a pie chart of the nodes' status within each group. Nodes (yellow in the illustration) are shown based on the group order with their current deployment state shown by the status light.

Nodes are automatically arranged into a series of default groups where each default group represents a switch. The default groups are named using the MAC address of the switch associated with them. Nodes are shown within each group ordered by port order (lowest on top). Switches and ports are discovered automatically during the configuration process. An italic group name indicates that the group is a default group, and a non-italicized group name indicates that the group is a user created group.



 If you use a consistent pattern for connecting nodes to switches then the Crowbar display matches your nodes' physical location

You can override the default group behavior from the Dashboard. To create a new group, enter the desired group name and click "Add Group," then drag the desired node into the empty group that

appears on the top left of the Dashboard. To rearrange nodes between groups, drag the node into the name of the desired group. To remove a node from a user-defined group setting, you must drag the node into the "[here]" portion of the hint text area.








 Nodes may be moved between user-defined groups, and between a user-defined group and a default group, but cannot be moved between default groups


Table 6: Deployment States

Status	Icon	Comment	User Action
Ready		Requested services provisioned.	Configure as needed.
Waiting	 blinking	Waiting for user input.	Node waiting to be allocated. See "Bulk Edit" or include node in a proposal.
Pending	 solid	Hardware and operating system installation.	<i>None: Crowbar is provisioning the node.</i>
In Process	 spinning	Crowbar and Chef actively deploying.	<i>None: Crowbar is provisioning the node.</i>
Failed	 blinking	Failure detected operating on node.	Correct issue.
Unknown		In between states or not reporting for 20 minutes (likely powered off).	Restart server if desired.


The Admin node is the node that runs Crowbar, Chef, and other core services. It exists in the system before the user interface is available.

Node Details

Clicking on a node's name displays details about the selected node in the details panel (green in the illustration). The details panel displays important information about the node, including its Full Qualified Domain Name (FQDN), uptime, switch connectivity hardware profile, and a detailed list of all active network connections.

 Node detail shows only a subset of the total details that Chef tracks for each node. To see the complete list, examine the Run List and Attributes for each node in Chef.

The Links list is barclamp-specific and expands depending on which barclamps are applied to the selected node. Links open a new window to view additional information about the node. The Barclamps and Roles lists indicate what capabilities have been assigned to the node.

 When a role or barclamp is selected in the details panel, the nodes that share the same barclamp or role are highlighted in the group panel. This helps quickly identify groups of nodes that are similarly configured

The buttons on the top of the details panel (Identify, Power On, Shutdown, and Reboot) use the node's IPMI interface to change the node's physical state. These states cause the node status to be unknown. These buttons are available only if the system is able to successfully configure the BMC on the target system.

Table 7: Details Panel Buttons

Button	Action	Useful When
Identify	Causes the identify light to blink for 15 seconds.	Trying to identify a node within a rack.
Power On	Sends a power-on signal to the BMC of the selected system.	Remotely powering on a system.

Button	Action	Useful When
Shutdown	Sends a power-off signal to the BMC of the selected system.	Remotely powering off a system.
Reboot	Sends a power-cycle signal to the BMC of the selected system.	Remotely power cycling a system, which has stopped responding.


The buttons on the bottom of the details panel (Delete, Reset, Reinstall, and Hardware Update) reset the node's deployment state. These functions are very useful during lab configurations when the system is being continuously reconfigured. The buttons take the following actions:

Table 8: Buttons on Bottom of Details Panel

Button	Action	Config Lost?	Reboot?	Useful When
Delete	Completely removes all records of the node from the Crowbar/Chef database. If a node is deleted, it is rediscovered when it reboots.	Yes	No	Removing nodes.
Reset	Removes all the roles assigned to the node and reimage it back to an unallocated node.	Yes	Yes	Reallocating the node for a new purpose.
Reinstall	Reimages the node and then reapplies the current deployment profile to the node. This rebuilds the server, returning it to a freshly deployed state.	No	Yes	Tuning the Chef recipes or configuration details.
Hardware Update	Keeps the current configuration, forces the node to reboot and apply BIOS and RAID updates.	No	Yes	Applying BIOS or RAID updates.


Using the Edit link (after the node name in the top left) lets you make per-node decisions about how the node is deployed.

To allocate a node, manually allocate the node from the edit page or include that node in an applied barclamp proposal.

 No changes are made to a node when it is discovered. This provides a safeguard against data loss if a system is accidentally PXE booted on the admin network. If this happens, delete the system from Crowbar, reboot it, and set the boot order in the BIOS to boot from the local disk first. Proceeding beyond the discovered state for a system by allocating it will immediately wipe all data on the system

Bulk Edit

The Bulk Edit screen allows you to quickly update the alias, description, RAID, BIOS, and allocation state for all the nodes in the system.

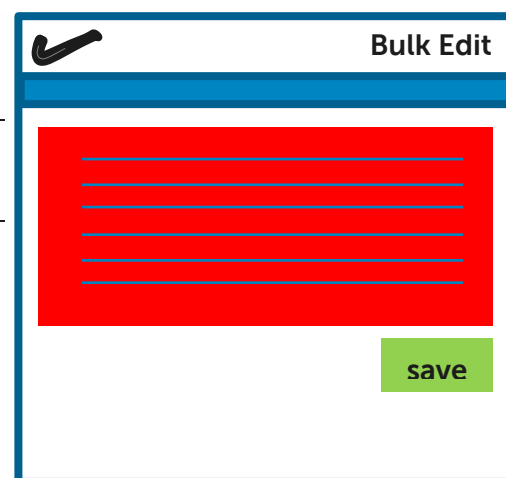
 Allocation is a destructive process. Allocating a node will immediately wipe all data on the node. You can return a node to the unallocated state using the Reset button as described above, though this will not restore any data.

The RAID and BIOS selections highlight the nodes' current value using [brackets]. The choices offered for BIOS and RAID selectors are determined by the BIOS and Dell RAID barclamps. They can be expanded after installation.

If no change is made then the node will not be updated.


Network

The Network menu exposes Crowbar network data and configuration options. Unlike the Nodes menu, the purpose of this menu is to visualize and manage the system network topology. Any barclamp is able to extend the Network menu.



Switches


The Switches menu item attempts to recreate the switch/port layout of your data center using information collected from the nodes. While the node view shows only the information about the node's administrative network, the Switch view shows what node interfaces are plugged in to which switch ports.


 If the 29-48 unit ports are not used or are inactive, then the switch view will show only a 24-port switch instead of a 48-port switch

The Switch view models either 24- or 48-port switches with the Switch MAC address and switch unit number at the top of the view. Ports are shown in two rows numbered from top-left to bottom-right with port numbers and include a status light that matches the nodes' current status. Inactive ports use the grey (●) status indicator and darker background color.

Hovering over any port causes two actions:

1. A tag showing the node's interface identifier (e.g.: eth1), alias, and description will be shown.
2. All other ports also connected to that node will be highlighted using a red cell border. This highlighting helps quickly identify a node's connection.

 If you want to just see the connections for a single node, you can add "?node=[node name]" to the URL. This call is linked on the switch name from the node view



The Switch view uses a compact layout to render switches. To expand the view to show additional details without hovering over the node, click on the switch name. This will reveal the expanded switch view. If the switch is 48 ports, then the expanded view will be split into 1..24 and 25..48 port blocks.

The expanded view shows the port number, status, network interface, and node alias inside the port cell. The same node highlight hovering effect is applied on the expanded view. Screen rendering may not be regular depending on the length of the node alias.

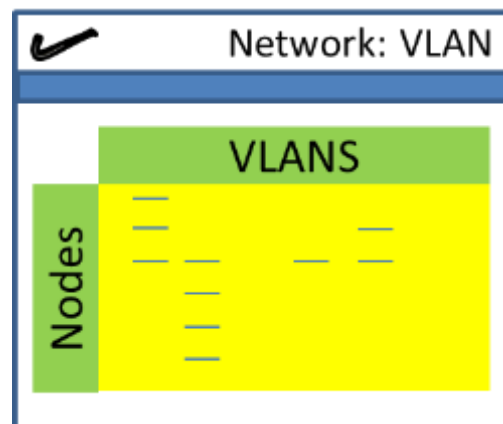
You cannot edit or change the switch configuration from this view.

VLANs

The VLANs menu item builds a matrix mapping between the VLANs tracked by the network barclamp and the nodes. This matrix represents both active and inactive connections between the nodes and the VLANs.

The VLAN view is a simple table with VLANs represented by columns and nodes as rows. If a node is active on a VLAN then "Active VLAN" is shown at the intersection. If a node is connected but inactive then "not active" or "N/A" is shown. If a node and VLAN are not mapped then the cell is empty. The IP address that the node has on a given VLAN can be seen by hovering the mouse over the appropriate cell. The VLAN ID can be seen by hovering the mouse over the network name.

You cannot edit or change the switch configuration from this view.



Barclamps

The Barclamps page lets you create, edit, review, and deploy proposals for barclamps. These activities are the way that Crowbar decides which nodes to deploy and how to configure them.

Barclamp List

The All Barclamps page shows a list of all available barclamps. The barclamps are represented as the blue lines in the figure to the right. Expand the barclamp, by clicking on the ▸ under the barclamp's name, to display the associated proposals for the selected barclamp, indicated by the red box in the figure. You jump directly to the relevant proposal by clicking on its name under its barclamp.

A barclamp will show the status of the proposals that are attached to it using a status light (see Table 9: Proposal Status below). If multiple proposals are assigned, then multiple lights are displayed. If there are no proposals, a diamond is displayed. Hovering over the light will show you the name and status of the matching proposal. The proposals status updates automatically without a refresh.





Table 9: Proposal Status

Status	Icon	Next Step	Comment	User Interaction
No Proposal	◇	Create	No proposal has assigned to the barclamp.	Create a proposal for the barclamp if desired.
User Input	●	Delete or Apply	Proposal waiting for user input and activation.	Edit the proposal. Apply proposal after review.
Active	●	Deactivate or Apply	Proposal has been deployed.	Ready for use.
Pending	● blinking	Wait or Dequeue	Queued for deployment.	Crowbar is building nodes during this phase.
In Progress	● spinning	None	Proposal is being configured.	None. Crowbar is working.
Failed	● blinking	Apply	Proposal failed during Apply	Correct error and reapply proposal.

From the Barclamp list, you may take actions on the proposals based on their state, as shown in the table above. Please review the *Barclamp Life Cycle* section for more information about the different proposal states.


All core barclamps automatically create proposals and do not allow users to create additional proposals. Some barclamps allow the creation of multiple proposals. These additional proposals can be used to manage deployment configurations or control which parts of the system are active in which barclamps.

From each barclamp, you can create a single new proposal for the system..

 Crowbar supports one proposal only per barclamp.

Proposal View/Edit

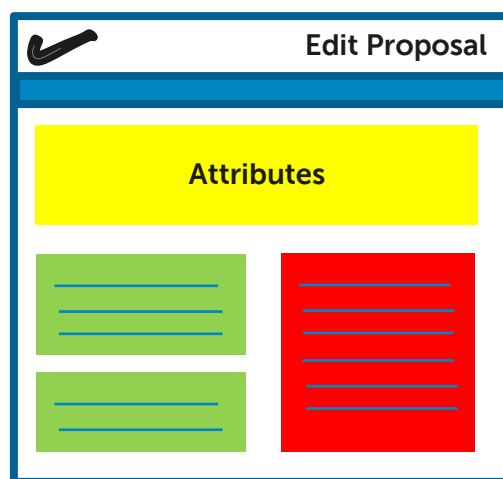
Selecting a proposal from on the list navigates to the proposal details page.


 If a proposal is active, you will be initially taken to a read-only view of the proposal. It is acceptable to edit an active proposal and re-apply. To access the proposal editor from the read-only view, click the edit button

Clicking the **Edit** button opens the Edit Proposal page. All proposals have two primary edit areas: Attributes (yellow in figure) and Node Deployment. Attributes are configurable data that is used by the Chef recipes. Node Deployment shows the Chef roles and nodes assigned to those roles.

Since each barclamp has unique attributes and roles you should consult the documentation for each barclamp if you plan to change its defaults.


Each barclamp may provide a custom editor for its attributes and node deployment information. The typical custom editor lets you set attribute values using a form, and drag and drop nodes from the available list (left column) into the roles associated with the barclamp (right, red on figure). Each barclamp may have specific logic that requires minimums or maximums for node assignments.




 While most barclamps coordinate with Chef to perform node deployments, Crowbar includes some special-function barclamps that can be used to change how Crowbar operates.

If the barclamp does not have a custom editor or your browser does not support the editor, Crowbar automatically uses a raw JSON editor. You can also use this view if you want to see the entire configuration details. Selecting the Raw view option on the right side of the Attributes or Deployment panel opens the JSON editor for that section. This option lets you directly edit the JSON configuration details for the proposal. This option is typically used when developing new barclamps or for advanced users only.


When you have finished editing the proposal, you may save or apply it. "Save" retains your configuration settings. "Apply" saves and then applies your proposal so that Crowbar begins deploying the barclamp on the selected nodes. Deleting a proposal removes it from the system and you lose your configuration.

 If you attempt to apply a proposal to nodes that have not yet been allocated, then Crowbar will queue the proposal while it automatically allocates the nodes. Following allocation, the proposal will be dequeued and applied to the nodes.

 When you apply a proposal, Crowbar creates Chef roles, and then puts them into the run list of the selected nodes. Crowbar uses a naming pattern for Roles that lets you quickly figure out which barclamp and proposal is being applied to a node's run list in Chef. The instantiated barclamp naming pattern is [barclamp]-config-[proposal]. Barclamps then use additional roles to control node proposal membership (aka the Run List).

Utilities

The Utilities menu is used by various barclamps to provide general-purpose tools that do not map within the normal node, network, or barclamp context. Any barclamp is able to extend the Utilities menu.

 The Utilities menu is a native part of the Crowbar barclamp; however, several menu items are populated by barclamps.

Exported Items

The Exported Items menu shows a list of files that have been prepared for export. There will be no items available until an export has been requested (see below). It also includes two button options:

- Log Export
- Chef Export

The exported items are shown grouped by category. Clicking on an item will download it through the browser download function. Clicking the delete icon (🗑️) will delete the file.

The export items page refreshes automatically.

Log Export

Log Export compiles all the system logs into a time-stamped single compressed archive. Typically, the first step in any troubleshooting activity will be to capture and send the logging information.

This action may take several minutes to complete and runs in the background. The Exported Items screen will automatically update when the export is complete.

Chef Export

Chef Export creates a time-stamped archive snapshot of the underlying Crowbar database. This snapshot can be used to replicate the installation for troubleshooting or backup.

This action generally takes only a few seconds to complete. Crowbar does this work in the background so that you can continue to use the UI while the export is operating.

Tempest

Tempest enables a set of automated integration tests that you can run against a live cluster.

To Run a Tempest Integration Test

1. Navigate to **Utilities > Tempest**.
2. Select the node from which you wish to run the report from the drop-down selector. Only nodes that have been allocated with the *tempest-server* role from the Tempest barclamp are available.
3. Click on the **Go!** Button to initiate the Tempest test run.

This action may take several minutes to complete. Crowbar does this work in the background so that you can continue to use the UI while the tests are running.

To View Tempest Test Results

1. Navigate to **Utilities > Tempest**.
2. Select the test results set by clicking on the **Ready** link.

Swift Dispersion

Swift Dispersion is a reporting tool to measure the overall health of a cluster.

To Run a Swift Dispersion Report


1. Navigate to **Utilities > Swift Dispersion**.
2. Select the node from which you wish to run the report from the drop-down selector. Only nodes that have been allocated with the *swift-dispersion* role from the Swift barclamp are available.
3. Click on the **Go!** Button to initiate the Swift Dispersion report.
4. Wait for the report status icon to turn green (*Ready*).

This action generally takes a short while to complete. Crowbar does this work in the background so that you can continue to use the UI while the report is generating.

To View Swift Dispersion Reports

1. Navigate to **Utilities > Swift Dispersion**.
2. Select the report by clicking on the **Ready** link.

The report displays in the web browser window.

 You can view the report without formatting by selecting the **Raw JSON** link.

Core Crowbar Barclamps

There are a set of core barclamps that are a part of every Crowbar installation. They are listed in the table below.

Table 10: Barclamps

Barclamp	Function / Comments
<u>Crowbar</u>	The roles and recipes to set up the barclamp framework. References other barclamps. Modify the default proposal to change the Usernames and passwords for access to the Crowbar UI.
<u>Deployer</u>	Initial classification system for the Crowbar environment (aka the state machine).
<u>Provisioner</u>	The roles and recipes to set up the provisioning server and a base environment for all nodes.
<u>Dell Branding</u>	Enables the Dell branding color palette in the Dell EULA version of Crowbar.
<u>IPMI</u>	Allows management of the IP Management Interface (IPMI) on servers when the BMC network is enabled.
<u>Network</u>	Instantiates network interfaces on the Crowbar managed systems. Also manages the address pool.
<u>Dell RAID</u>	Sets up LSI® RAID controllers in a variety of configurations. If missing, the RAID controllers can be set up manually.
<u>BIOS</u>	Configures BIOS options for Dell™ PowerEdge™ servers. If missing, the BIOS options can be configured manually.
<u>DNS</u>	Manages the DNS subsystem for the cluster.
<u>Logging</u>	Centralized logging system based on syslog.
<u>Git</u>	Provides a Git instance.
<u>NTP</u>	Common NTP service for the cluster (required for secure access). An NTP server can be specified.
<u>Nagios</u>	System monitoring service for the cluster that can be used by other barclamps.
<u>Test</u>	Provides a shell against which you can write tests.
<u>Ganglia</u>	Performance monitoring service for the cluster that can be used by other barclamps.

Details about these barclamps are provided below.

Crowbar Barclamp

The Crowbar barclamp provides the roles and recipes to set up the barclamp framework.

The Crowbar barclamp initializes the system, creates initial proposals of other barclamps defined in its configuration, and creates the users to access the Crowbar API and UI. By default, the system creates Network, Ganglia, Nagios®, NTP, DNS, Provisioner, Deployer, IPMI, RAID, and BIOS proposals based upon the default configuration of their barclamps. The initialization function of the Crowbar barclamp works exactly like other barclamps. A proposal is created and can be committed during installation.

All barclamps' transition functions can be called directly, but the Crowbar barclamp calls these in an order specified in its configuration, which is determined by their priority. The default unspecified priority is 100. The special cases are the Provisioner, which is last, and the Deployer and Network, which are first and second, respectively.

Table 11: Crowbar Barclamp Parameters

Name	Default	Description
instances	The starting barclamps using their default configurations.	A map of barclamp names that reference a list of JSON files (default is special to mean to take the defaults) that represent barclamp instances to create.
users	A map of users — containing Crowbar.	This map defines the users allowed to access Crowbar's UI and REST API.

The users map contains a map. The key is the user name and the rest of the required fields are as follows:

Table 12: User Name Key

Name	Description
password	Clear text password of the user.
description	A description of the user.

Deployer Barclamp

The Deployer provides an initial classification system for the Crowbar environment. As nodes are discovered, the Deployer makes sure that discovery tools are run on the node by making sure that the Deployer-client role is assigned to the node. The results of that discovery are classified, and the node's attributes are updated to reflect its potential usage. The Deployer also builds a map of valid and usable disks.

The Deployer gives the primary name to the node at the discovered state. The names default to the letter "d" and the MAC address (with dashes instead of colons). The Deployer also allocates the admin and BMC addresses from the network barclamp.

In addition, the Deployer defines and provides the node's configuration for RAID and BIOS. These values are assigned part of the hardware-installing state transition. The Deployer uses a list of role name patterns that define what the RAID and BIOS configurations should be. These are applied as values in the node attributes under crowbar -> hardware. bios_set can be either *Virtualization* or *Storage*. RAID set can be either *JBODOnly* or *SingleRaid10*.

The Deployer is also responsible for manipulating the run list during the hardware-installing and update (or hardware-updating) states. The run list should only include BIOS, RAID, and IPMI operations.

The Deployer also controls the allocate flag on the node. The allocate flag is used to pause the node after discovery. The node waits for it to be allocated to continue. The Deployer has a configuration option to indicate if the allocate flag should be set to false (and cause a pause) or just allocate all nodes.

Table 13: Deployer Barclamp Parameters

Name	Default	Description
bios_map	A list of default settings for bios and raid for swift and nova.	The map defines a list of patterns that would apply a configuration setting for BIOS and RAID.
use_allocate	true	A Boolean value true indicates that a pause should be injected after the discovered state to allow the admin to accept and allocate the node.

Table 14: BIOS Map Entry Keys


Name	Description
pattern	Regular expression applied to the role names on the node.
bios_set	The BIOS set of parameters to apply. Values are: Virtualization or Storage.
raid_set	The RAID set of parameters to apply. Values are: JBODOnly or SingeRaid10.

Provisioner Barclamp

The Provisioner provides the roles and recipes to set up the provisioning server and a base environment for all provisioned nodes. The Provisioner also provides the transition entry point for nodes that need to have DHCP transitions done. The Provisioner assumes that IP addressing is handled outside of this barclamp.

Table 15: Provisioner Barclamp Parameters

Name	Default	Description
default_user	crowbar	User to create for external login.
default_password	unset	Clear text password to use for external login.
default_password_hash	Hash of crowbar	MD5 hash of password to use for external login. <i>printf 'password' mkpasswd -s -m md5</i> will generate the hash.
web_port	8091	The default Web port that the repository web server uses.
use_local_security	true	This defaults the security updates path in the install to use the admin node instead of the Internet.
dhcp	map	This is a map that contains the DHCP parameters (lease-time and state machine).
lease-time	60	The number of seconds a DHCP lease is valid for the system.
state_machine	map	This is the state machine that the DHCP server uses in this instance of the barclamp.

 While neither is required, one of *default_password* or *default_password_hash* is required.

Network Barclamp

The Network barclamp provides two functions for the system. The first is a common role to instantiate network interfaces on the Crowbar managed systems. The other function is address pool management.

The network interfaces are controlled by the network role that is applied by the barclamp as a node transition to "installed". Based upon assigned addresses, the network recipe creates the appropriate single, dual, or team mode interface sets.

The network assignment function is handled by the creation of an API extension of the base barclamp. The barclamp adds the `allocate_ip` REST API call. This function allocates an IP address from a requested network and updates the node's attributes and the network's data bag. The available networks (and their parameters) are defined in the configuration for the barclamp.

Modification of the following parameters should only be done when installing Crowbar, prior to running the `./install systemname.yourdomain.com` command. See the *Crowbar OpenStack Deployment Guide* for more information.

Table 16: Network Configuration Options

Name	Default	Description
mode	single	A string value of single, dual, or team. This specifies the default network interface construction model.
teaming	map	A map of values specific to teaming.
networks	map	A map of networks that this barclamp should manage.

Table 17: Teaming Sub-Parameters

Name	Default	Description
mode	6	The default teaming algorithm to use for the bonding driver in Linux.

Table 18: Default Networks

Name	Usage	Notes
admin	Private network for node-to-node communication	A router, if wanted, is external to the system. This network must be owned by the Crowbar system to run DHCP.
bmc	Private network for BMC communication	This can be the same as the admin network by using the ranges to limit what IP goes where. A router, if wanted, is external to the system.
bmc_vlan	Private network for admin nodes on the BMC network	This must be the same as the BMC network and have the same VLAN. This is used to generate a VLAN-tagged interface on the admin nodes that can access the BMC LAN.
storage	Private network for storage traffic	A router, if wanted, is external to the system.
public	Public network for Crowbar and other components	A router, if wanted, is external to the system.

Table 19: Network Parameters

Name	Default	Description
vlan	Integer	The VLAN to use on the switch and interfaces for this network.
use_vlan	true	A value of true indicates that the VLAN should apply to the interface. A value of false assumes that the node receives untagged traffic for this network.
add_bridge	false	Indicates if the network should have a bridge built on top of it. The bridge will be br. This is mostly for Nova compute.
subnet	IP Address	The subnet for this network.
netmask	Netmask	The netmask for this network.
router	IP Address	The default router for this network.
broadcast	IP Address	The default broadcast address for this network.
ranges	map	This contains a map of strings to start and stop values for the network. This allows allocating sub-ranges with the network for specific uses. For example: DHCP, admin, BMC, hosts.

Table 20: Range Map String Key

Name	Type	Description
start	IP Address	First address in the range, inclusive.
end	IP Address	Last address in the range, inclusive.



Settings in the Network barclamp should not be changed after Admin Node installation.

Dell RAID Barclamp

A RAID (Redundant Array of Independent Disks) controller can make multiple disks look like one big, smart, safe disk. Different RAID controllers have different capabilities. In general, a RAID controller can support one or more RAID volumes of various types. Any disks not included in a RAID volume are directly exposed to the operating system. Disks that are directly exposed like this are known as Just a Bunch Of Disks (JBOD).

Operations

The Dell RAID barclamp is responsible for configuring the storage subsystem as directed by Crowbar. Users have the option of manually choosing the RAID configuration to use. The Dell RAID barclamp can be enabled or disabled by editing the default RAID proposal. The Dell RAID barclamp currently supports specific RAID controllers as defined in the Reference Architecture.

The Crowbar code makes sure that the configuration on the RAID controller matches that specified within the Crowbar configuration.

The parts that determine the configuration for a node are:

- A set of Chef data bags, which contain the RAID configuration (in data bags/crowbar-data). The defaults are SingleRaid10 and JBOD Only.
- Crowbar (the Deployer barclamp) sets the above property when a node is allocated to a proposal.

When invoked, the recipe uses a Chef Lightweight Resources and Providers (LWRP) to inspect the current configuration on the system and compare it to the desired state. If the two diverge, the code will:

- Delete any RAID sets that are no longer required.
- Allocate available disks among the desired RAID sets, according to the order attribute.
- Issue commands to apply the configuration.

The following RAID configurations are available:

- **JBOD** - All disks are independently addressed, with no collective properties. There is no minimum to the number of disks required in this configuration. On controllers which do not natively support JBOD, this selection is simulated by creating a Single Disk volume RAID0 for each physical disk.
- **RAID 10** - Sometimes called RAID 1&0 or RAID 1+0. This configuration is a stripe of mirrors and requires at least 4 drives to be implemented. If an insufficient number of drives is present, this selection is automatically replaced with a RAID1 configuration (which requires a minimum of 2 drives).

The Dell RAID barclamp only changes the controller's configuration when a node is initially installed by Crowbar (i.e., on first install, or after the node has been reset through Crowbar). This prevents accidental loss of all data when performing a hardware-update operation.

Supported RAID Controllers

Crowbar currently supports a few models of raid controllers, based on LSI® technology:

Table 21: Supported RAID Controllers

Controller	Platform	Notes
9202	C8000	sas2ircu supported
2008	C8000/C2100/C6100/C6105/C6220	sas2ircu supported
9265	C6220	MegaCli supported
9260	C2100/C6100	MegaCli supported
H710	R720/R720xd	wsman supported
H310	R720/R720xd	wsman supported

Each controller has different capabilities:

Table 22: RAID Controller Capabilities

Controller	Capabilities
9202	Non raid capable: disks are exposed to the operating system (JBOD)
2008	Up to two (2) separate RAID volumes of either RAID0, RAID1, RAID1E, or RAID10. Any disks not included in a RAID volume are directly exposed to the Operating System (i.e., JBOD).
9265	Up to 32 volumes, with up to eight (8) spans each can be configured. A span can be a single disk or a two-disk mirror. Volumes are identified by numbers; no name can be assigned. JBOD is only supported as a single RAID0 volume.
9260	Up to 32 volumes, with up to eight (8) spans each can be configured. A span can be a single disk or a two-disk mirror. Volumes are identified by numbers; no name can be assigned. JBOD is only supported as a single RAID0 volume.
H710	RAID 10: uses single-member RAID0 arrays for JBOD configuration
H310	RAID 10: uses single-member RAID0 arrays for JBOD configuration

BIOS Barclamp

The BIOS barclamp provides the following specific control features for certain Dell servers:

PowerEdge-C Series


- Updates the BIOS image to a known, consistent version depending on the hardware
- Updates the BMC image to one compatible with the BIOS version
- Configures the BIOS parameters to match the role assigned to the node

PowerEdge-R Series

- All updateable components are updated via the *wsman* interface:
 - BIOS
 - Lifecycle controller firmware
 - RAID controller firmware
 - NIC firmware
 - And others

The BIOS barclamp supports specific hardware platforms as defined in the Reference Architecture document. These operations are performed when the node is first allocated to use, or when the hardware-update operation is selected from the Crowbar user interface.

All of the BIOS barclamp's responsibilities can be enabled or disabled individually, by editing the default BIOS proposal.

 Dell highly recommends that you either enable or disable the BIOS Update and BMC update functions together.

Updating the BMC image can take up to 30 minutes, and involves multiple system reboots during which the BMC will be inaccessible.

Operations

In its current state, the BIOS barclamp supports the following Crowbar integration, CentOS, and Ubuntu capabilities:

Crowbar Integration

- When a node transitions to the discovered state it will have the **bios** role added to it
- When a node is allocated (manually or via committing a proposal)
 - The barclamp checks to see if any firmware can be updated. How this is done depends on system type:
 - On PowerEdge-R (PE-R) series systems, we will try to update the system to the current firmware versions via *wsman*.


- On PowerEdge-C (PE-C) series systems, we will check to make sure that the BIOS and BMC are at the minimum specified versions for each platform, and attempt to update them if they are not.
- The barclamp attempts to update firmware a maximum of three (3) times before giving up, and placing the node in a problem state. This forces a reboot and retry loop until the issue is fixed, and the update count is modified to less than three.
 - If the update issue is not fixed (flashing updates fails multiple times, and sets the count of attempts to three), the node will remain in a problem state even after completion of all other hardware configuration (BIOS parameters, RAID configuration, etc.).
- The `*bios_set*` parameter the barclamp will use to choose the parameter set.
 - On C6100, C6015, and C2100 systems the barclamp only tries to set BIOS parameters if the system is in the version whitelist contained in `legacy_token_map_maker.rb`. This ensures that the updated tokens are valid for the BIOS revision that the system is running on.

Sledgehammer

Sledgehammer is the hardware discovery and configuration component of Crowbar. It is the tftboot image that PXE clients will load when they network boot. It is a modified version of CentOS that reports discovery information back to the Crowbar server.

Because Sledgehammer is CentOS based, Dell Update Packages (DUPS) run in the context of Sledgehammer:

- Check if BMC needs updating, and update if needed
 - Reboot if updated
- Check if BIOS firmware needs updating
 - Reboot if updated
- Update BIOS parameters

 DUPS do not work on Ubuntu, which is why Crowbar uses Sledgehammer in CentOS.

Firmware Updates

Table 23: Supported Firmware Updates


Platform	Updated BIOS Version	Updated BMC Version
C2100	01/27/13 B23	09/19/12 1.82, A01
C6100	11/1/2011 1.64, A00-1.64	11/1/2011 1.24, A08-1.24
C6105	03/20/13 2.2.0	03/20/13 1.30
C6145	01/23/13 3.0.0	01/23/13 1.11
C6220	05/30/13 1.1.19	05/30/13 2.02
C8220	03/20/13 1.1.17	03/20/13 2.01

Updating BIOS and BMC Firmware

1. Download the latest BIOS update files from the Dell Support Site (<http://www.dell.com/support/drivers>).
 - **PE-C Series platforms** - BIN files
 - **PE-R Series platforms** - DUP files (Windows .exe format)
2. Using your tool of choice copy the files to the platform-specific directory in the Crowbar admin node's `/updates` directory.
 - For example, for a C6220, copy to `/updates/6220`.
3. For PE-C Series platforms:
 - a. Click on the **Barclamps** tab, and then select **All Barclamps**.
 - b. Select the **BIOS** barclamp.

- c. In the *Edit Proposal* screen, select **Bios: Default > Raw**.
- d. Update the BIOS file names to the names of the files you uploaded in Step 2 above.
- e. Click on the **Save** button, and then click on the **Apply** button.
4. For PE-R Series platforms:
 - Go to Step 5.
5. Navigate to **Nodes > Dashboard**, and then click on the **node name**.
6. Click on the **Hardware Update** button to apply the updates to a deployed node.

Whitelisting New BIOS Versions

 Be **very** careful with this procedure. Indiscriminate use of the instructions below can result in non-operational nodes.

Occasionally some PE-C Series servers may ship with BIOS versions newer than that which Crowbar currently supports. Crowbar will not configure BIOS parameters if the version flashed on the system is newer or not recognized. In addition, if the whitelist is amended manually, setting BIOS parameters through Crowbar may result in non-operational nodes.

If this occurs, you can whitelist your PE-C server's updated BIOS version. This will ensure that Crowbar is aware of the node's BIOS version, and will allow the boot sequence to continue.

To Whitelist an Updated BIOS Version:

1. Shut down the failed node.
2. Initiate an SSH session to the Crowbar admin node (typically 192.168.124.10, using port 22). For example:

```
ssh 192.168.124.10:22
```

3. Login as user crowbar (default password crowbar), and then switch user to root. For example:

```
sudo -i
```

- a. Enter the password, `crowbar`.
4. Edit the node's hardware-specific data bag with your editor of choice. For example:


```
less /opt/dell/barclamps/bios/chef/data_bags/crowbar-data/bios-set-PowerEdgeC6220-default.json
```
5. Edit the "versions" declaration to include the BIOS version for your hardware. For example:

```
"versions": ["1.0.21", "1.0.28", "1.1.19"],
```

 Ensure that within the brackets, you enclose each BIOS version within quotation marks, and append a trailing comma.

6. Save the data bag file, and restart the failed node.
7. Reinstall the BIOS barclamp:

```
/opt/dell/bin/barclamp_install.rb bios
```

Crowbar will recognize changed BIOS defaults, and perform BIOS configuration as appropriate.

Dell Branding Barclamp

This barclamp is enabled by default in the Dell EULA version of Crowbar. It provides a different color palette than the open-source version.

 The Dell Branding barclamp does not appear in the *Crowbar Members* screen's list of core Crowbar barclamps.

IPMI Barclamp

The IPMI barclamp configures IPMI access on platforms that support it. It selects an IP address and assigns it to the BMC. Prior to doing this, it checks the IP address that is currently assigned to the BMC. If the address falls within the range of addresses that are configured for the BMC network in the network JSON, then it will use that IP address instead of assigning a new one.

The IPMI barclamp allows the user to control the following parameters:

Table 24: IPMI Barclamp Parameters


Name	Default	Description
bmc_enable	true	Controls if the barclamp attempts to work on the BMC.
bmc_password	cr0wBar!	The password that will be configured on the BMC.
bmc_user	root	The username that will be configured on the BMC.
debug	true	Turns on more verbose output.

NTP Barclamp

The NTP barclamp provides a common NTP service for the cluster. You can specify an NTP server or servers and all other nodes are clients of them. By default, the time on all nodes are synced to the hardware clock of the Crowbar admin node. The time zone of all Crowbar managed nodes in the cluster is set to UTC.

Table 25: NTP Barclamp Parameters

Name	Default	Description
external_servers	empty list	A list of IP addresses or hostnames that should be used as external NTP servers. Hostname can be used if the DNS barclamp is configured to have access to an external resolver.
admin_ip_eval	"Chef::Recipe::Barclamp::Inventory.get_net_work_by_type(node, \"admin\").address"	The Ruby eval expression that returns the admin IP address of a node.

 If you are setting up an external server it can take up to 5 minutes for the nodes to sync with the server. Systems should not be rebooted during this process. If they are rebooted, then they will pause during boot-up for time synchronization.

DNS Barclamp

The DNS barclamp provides a centralized management subsystem for the cluster.

Table 26: DNS Barclamp Parameters

Name	Default	Description
contact	support@<your_domain>	The administrative email address for your deployed cluster.
domain	<your_domain>	The organization specified upon Crowbar admin node installation.
forwarders	N/A	N/A
static	N/A	N/A

Logging Barclamp

The Logging barclamp provides a centralized logging system based on syslog. The barclamp enables a centralized log server that can then forward information to external syslog servers. The Crowbar installation process sends logs to the admin node by default, but the configuration from the logging barclamp can override this initial configuration.

Table 27: Logging Barclamp Configuration Parameters

Name	Default	Description
external_servers	Empty list	A list of IP addresses for the logging-server to which to forward logs.

Nagios® Barclamp

The Nagios® barclamp provides a common Nagios® service for the cluster. A Nagios® server or servers can be specified and all other nodes are clients of them. The barclamp attempts to direct all traffic over the admin network.

Table 28: Nagios® Barclamp Parameters

Name	Default	Description
admin_interface_eval	<code>Chef::Recipe::Barclamp::Inventory.get_network_by_type(node, \"admin\").interface</code>	The Ruby eval expression that returns the admin interface of the node.
admin_ip_eval	<code>\"Chef::Recipe::Barclamp::Inventory.get_network_by_type(node, \"admin\").address\"</code>	The Ruby eval expression that returns the admin IP address of the node.

Table 29: Ganglia Barclamp Parameters.

Name	Default	Description
interface_eval	<code>Chef::Recipe::Barclamp::Inventory.get_network_by_type(node, \"admin\").interface</code>	The Ruby evaluation string that gets the interface of the admin interface.

Test Barclamp

The Test barclamp provides a shell for writing tests against. It allows for failures to be injected and other barclamps can be validated against it.

This barclamp is currently used only to perform a quick smoke test of Crowbar following Crowbar installation, and as a result should never be deployed.

Table 30: Test Barclamp Parameters.

Name	Description
barclamps	A list of supported barclamps that are used as the return value for the barclamp list API call
instances	A map of barclamp names that reference a list of JSON files (default is special to mean to take the defaults) that represent starting barclamp instances to create

Ganglia Barclamp

The Ganglia barclamp provides a common Ganglia service for the cluster. Ganglia server or servers can be specified and all other nodes are clients of them.

Table 31: Ganglia Barclamp Parameters

Name	Default
interface_eval	<code>Chef::Recipe::Barclamp::Inventory.get_network_by_type(node, \"admin\").interface</code>

Git Barclamp

The Git barclamp provides a Git instance to enable Pull from Source, a mechanism which installs the latest updates to barclamp packages from GitHub, at barclamp installation time.

Table 32: Git Barclamp Parameters

Name	Default	Description
update_origins	false	Enables/disables Pull from Source

Supplemental Material

System Verification

As a final step, it is important to verify that your deployment has succeeded. Crowbar provides limited feedback to confirm that the Chef recipes were successfully deployed.

You should consult the *Getting Started Guide* and barclamps specific to your system for details on verification of deployment.


Managing Barclamps

This section briefly describes barclamps, and how to import barclamps.

 For information about creating barclamps, please visit <https://github.com/dellcloudedge/crowbar/wiki/Barclamp:-create-&-install-steps>.

Introduction


A barclamp is a deployment module that is imported from its own code repository into the Crowbar framework. A barclamp cannot operate without Crowbar, but you do not have to create a unique build of Crowbar in order to create a barclamp.

 You must install Crowbar before importing barclamps.

Importing a Barclamp

Once you have created a barclamp, you can import the barclamp into Crowbar and Chef. Assuming that you already created the foo barclamp in /barclamps (see Creating a Barclamp), proceed as follows:

1. From the Crowbar server, become the super admin: `sudo -i`
2. Run the barclamp install script: `/opt/dell/bin/barclamp_install /barclamps/foo.tar.gz`
 - a. `"/barclamps/foo.tar.gz"` is the file name of your barclamp; it can be any name you wish.

 The core barclamps are in `/opt/dell/barclamps`.

3. Restart the Crowbar web server: `bluepill crowbar-webserver restart`

Your barclamp should now appear in the Crowbar UI. You can also see it in Chef under the Crowbar data bag.

While barclamps are generally safe to install multiple times, you can uninstall a barclamp using `"barclamp_uninstall.rb /path/to/barclamp"`.

More Information

Information on how to develop your own barclamp is available in the *Developers Guide*. If you are interested in creating, extending or contributing barclamps, please use one of the following contact methods.


Support

Crowbar Support

To obtain support for Crowbar:

- See the Crowbar wiki on GitHub: <https://github.com/crowbar/crowbar/wiki>.
- Gather log information (see below).
- Email the Crowbar listserv — join at: <https://lists.us.dell.com/mailman/listinfo/crowbar>.

To help facilitate troubleshooting of the environment, a utility to gather logs has been provided. Use the Log Export function in the user interface as detailed above, or browse to http://<Admin_ip>:3000/support/logs. This creates a tar archive of the relevant logs and asks the user for a location to save the resulting archive.

 Depending on the size of the logs to be gathered, this utility may take a while to run.

To Learn More

For more information on Crowbar, visit:

www.dell.com/crowbar

©2013 Dell Inc. All rights reserved. Trademarks and trade names may be used in this document to refer to either the entities claiming the marks and names or their products. Specifications are correct at date of publication but are subject to availability or change without notice at any time. Dell and its affiliates cannot be responsible for errors or omissions in typography or photography. Dell's Terms and Conditions of Sales and Service apply and are available on request. Dell service offerings do not affect consumer's statutory rights.

Dell, the DELL logo, and the DELL badge, PowerConnect, and PowerVault are trademarks of Dell Inc.